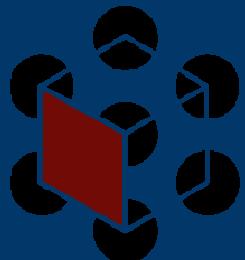


# Glift: Generic, Efficient Random-Access GPU Data Structures

Aaron Lefohn

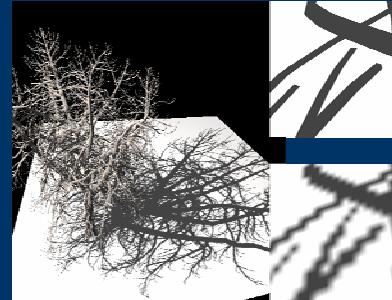
University of California, Davis



# Problem Statement

---

- “A data structure abstraction for graphics processing units (GPUs) can simplify the description of new and existing data structures, stimulate development of complex GPU algorithms, and perform equivalently to hand-coded implementations.”

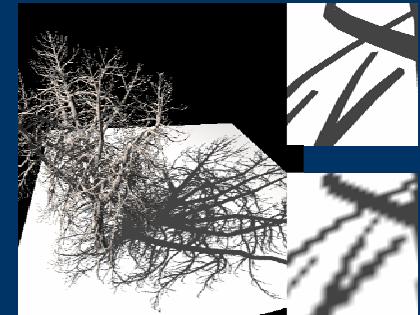


# Problem Statement

---

- **Goal**

- Simplify creation and use of random-access GPU data structures for graphics and GPGPU programming



- **Contributions**

- Abstraction for GPU data structures
- Glift template library
- Iterator computation model for GPUs



# Collaborators

---

- Joe Kniss  
University of Utah
- Robert Strzodka  
Stanford University
- Shubhabrata Sengupta  
University of California, Davis
- John Owens  
University of California, Davis



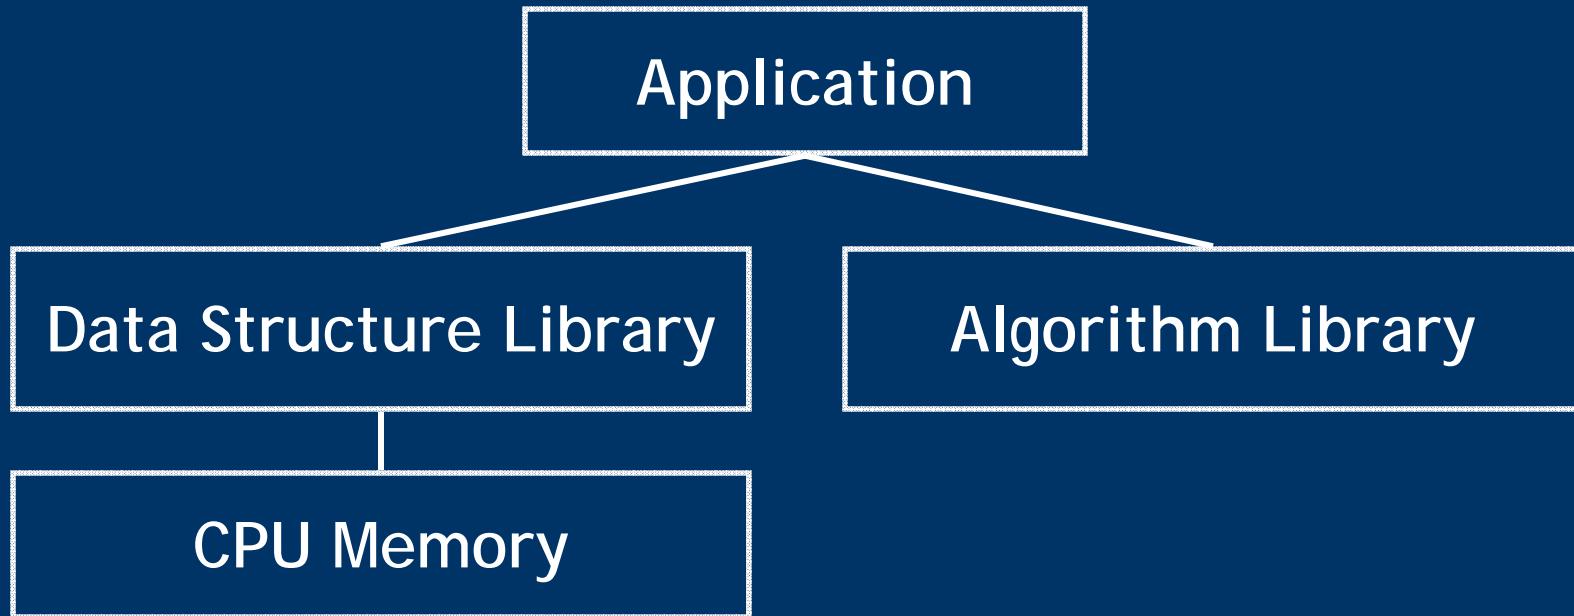
Aaron Lefohn  
University of California, Davis

# Many Interesting GPU Data Structures

- Photon map Purcell
  - Sparse matrix Boltz, Krueger
  - Sparse simulation grid Lefohn
  - Polycube (3D grid, cubeMap, ...) Tarini
  - N-tree Lefebvre
- 
- But...
    - No way to distribute/reuse implementations
    - Complexity stifles innovation



# CPU Software Development



- Benefits
  - Algorithms and data structures expressed in problem domain
  - Decouple algorithms and data structures
  - Code reuse



# GPU Software Development

---

Application  
- Data structure and algorithm

GPU Memory

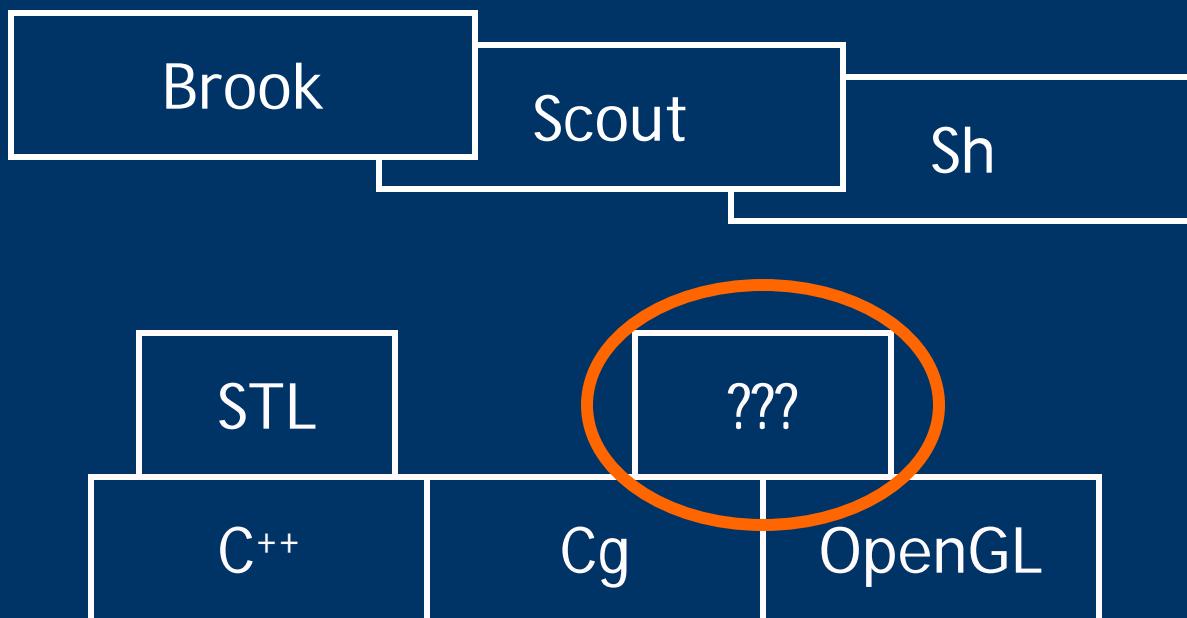
- Problems
  - Code is tangled mess of algorithm and data structure access
  - Algorithms expressed in GPU memory domain
  - No code reuse



# GPU Data Structures

- What's Missing?

- Standalone abstraction for GPU data structures for graphics or GPGPU programming



# Simple Example

- CPU (C++)

```
float srcData[10][10][10];
float dstData[10][10][10];

... initialize data ...

for (size_t z = 1; z < 10; ++z) {
    for (size_t y = 1; z < 10; ++y) {
        for (size_t x = 1; z < 10; ++x) {
            dst[z][y][x] = log( 1 + src[z][y][x] );
        }
    }
}
```



# We Want To Transform This...

- GPU (Cg)

```
float3 getAddr3D( float2 winPos, float2 winSize, float3 sizeConst3D ) {
    float3 curAddr3D;
    float2 winPosInt = floor(winPos);
    float addr1D = winPosInt.y * winSize.x + winPosInt.x;
    addr3D.z = floor( addr1D / sizeConst3D.z );
    addr1D -= addr3D.z * sizeConst3D.z;
    addr3D.y = floor( addr1D / sizeConst3D.y );
    addr3D.x = addr1D - addr3D.y * sizeConst3D.y;
    return addr3D;
}

float3 logAlg(uniform samplerRECT data,
               uniform float2      winSize,
               uniform float3      sizeConst3D,
               float2              winPos : WPOS ) : COLOR
{
    float3 addr3D = getAddr3D( winPos, winSize, sizeConst3D );
    float data   = texRECT(data, addr3D );
    return log( 1 + data );
}
```



# We Want To Transform This...

- GPU (Cg and C++)

```
float3 getAddr3D( float2 winPos, float2 winSize, float3 sizeConst3D ) {  
  
    float3 curAddr3D;  
    float2 winPosInt = floor(winPos);  
    float addr1D = winPosInt.y * winSize.x + winPosInt.x;  
    addr3D.z = floor( addr1D / sizeConst3D.z );  
    addr1D -= addr3D.z * sizeConst3D.z;  
    addr3D.y = floor( addr1D / sizeConst3D.y );  
    addr3D.x = addr1D - addr3D.y * sizeConst3D.y;  
    return addr3D;  
}  
  
float3 logAlg(uniform samplerRECT data,  
              uniform float2      winSize,  
              uniform float3      sizeConst3D,  
              float2      winPos : WPOS) : COLOR  
{  
    float3 addr3D = getAddr3D( winPos, winSize, sizeConst3D );  
    float data   = texRECT(data, addr3D );  
    return log( 1 + data );  
}
```

```
GLuint srcDataId = 1;  
glBindTexture(GL_TEXTURE_RECTANGLE_ARB, srcDataId);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_MIN_FILTER, GL_NEAREST);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_MAG_FILTER, GL_NEAREST);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_WRAP_S, GL_CLAMP);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_WRAP_T, GL_CLAMP);  
glTexImage2D(GL_TEXTURE_RECTANGLE_ARB, 0, GL_LUMINANCE32F_ARB,  
            0, 0, 40, 40, GL_LUMINANCE, NULL);  
  
GLuint dstDataId = 2;  
glBindTexture(GL_TEXTURE_RECTANGLE_ARB, dstDataId);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_MIN_FILTER, GL_NEAREST);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_MAG_FILTER, GL_NEAREST);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_WRAP_S, GL_CLAMP);  
glTexParameteri(GL_TEXTURE_RECTANGLE_ARB, GL_WRAP_T, GL_CLAMP);  
glTexImage2D(GL_TEXTURE_RECTANGLE_ARB, 0, GL_LUMINANCE32F_ARB,  
            0, 0, 40, 40, GL_LUMINANCE, NULL);  
  
... Initialize data ...
```



# Into This.

- GPU (C++ and Cg with Glift)

```
typedef glift::ArrayGpu<vec3i, vec1f> ArrayType;  
ArrayType src( vec3i(10,10,10) );  
ArrayType dst( vec3i(10,10,10) );  
  
... initialize data ...
```

```
float logAlg( ElementIter srcData ) : COLOR  
{  
    return log( 1 + srcData.value() );  
}
```



# Overview

---

- Motivation and Previous Work
- Abstraction
- Implementation
- Examples
- Conclusions



# Abstraction Design Goals

---

- GPU data structure abstraction that
  - Enables easy creation of new structures
  - Is minimal abstraction of GPU memory model
  - Separates data structures and algorithms
  - Encourages efficiency



# Building the Abstraction

---

- Approach
  - Bottom-up, working towards STL-like syntax
  - Identify common patterns in GPU papers and code
  - Inspired by
    - STL, Boost, Brook, STAPL, Stepanov



# What is the GPU Memory Model?

- CPU interface

|                                     |            |                |
|-------------------------------------|------------|----------------|
| • <code>glTexImage</code>           | malloc     |                |
| • <code>glDeleteTextures</code>     | free       |                |
| • <code>glTexSubImage</code>        | memcpy     | GPU -> CPU     |
| • <code>glGetTexSubImage*</code>    | memcpy     | CPU -> GPU     |
| • <code>glCopyTexSubImage</code>    | memcpy     | GPU -> GPU     |
| • <code>glBindTexture</code>        | read-only  | parameter bind |
| • <code>glFramebufferTexture</code> | write-only | parameter bind |

\* Does not exist. Emulate with `glReadPixels`



# What is the GPU Memory Model?

- GPU Interface (shown in Cg)

- `uniform samplerND`
- `texND(tex, addr)`
- `varying floatN stream`
- `stream`

data structure param declaration  
random-access read

stream parameter declaration  
stream read



# GPU Data Structure Abstraction

---

- Factor GPU data structures into
  - Physical memory
  - Virtual memory
  - Address translator
  - Iterators



# Physical Memory

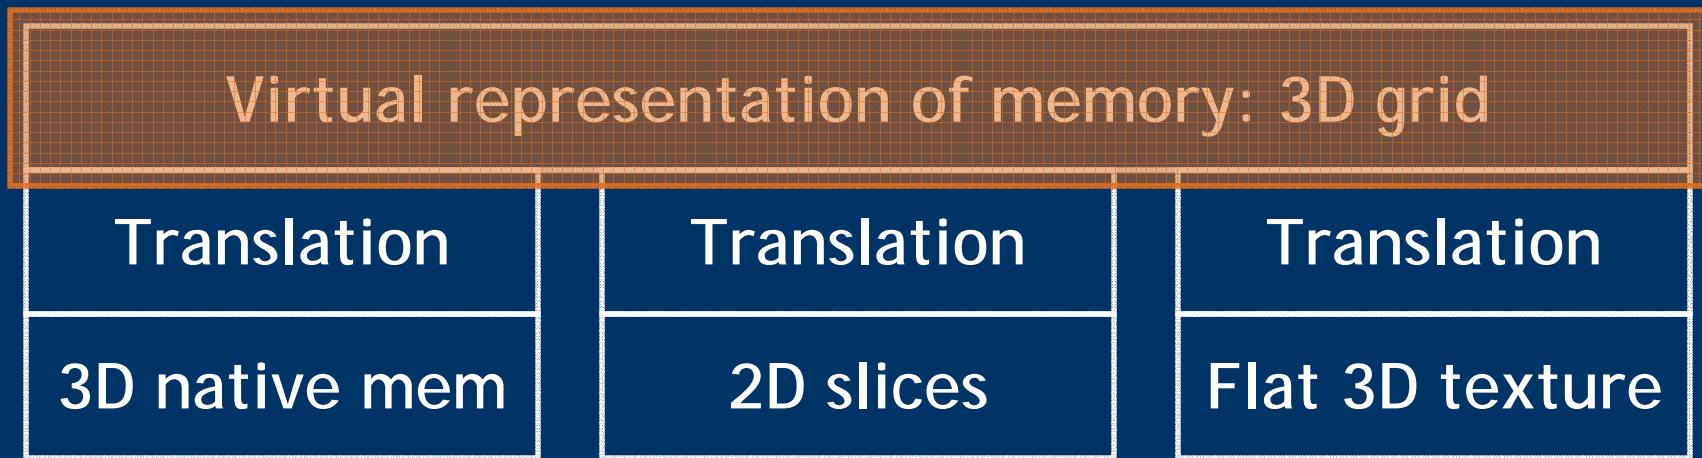
---

- Native GPU textures
  - Choose based on algorithm efficiency requirements
  - 1D, 2D, 3D, Cube, Mip
    - Dimensionality
    - Read-only vs. read-write
    - Point-sample vs. filtering
    - Maximum size



# Virtual Memory

- Virtual N-D address space
  - Choose based on problem space of algorithm
  - Defined by physical memory and address translator



# Address Translator

- Mapping between physical and virtual addrs



- Core of data structure
- Small amount of code defines *all* required CPU and GPU memory interfaces



# Address Translator

---

- Core of data structure
  - Extension point for creating new structures
  - Must define

`translate(...)`

`translate_range(...)`



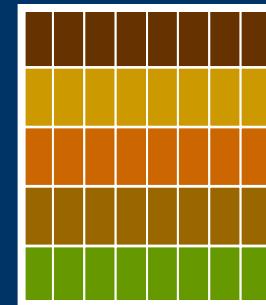
# Address Translator Classifications

- Representation
  - Analytic / Discrete
- Memory Complexity
  - $O(1)$ ,  $O(\log N)$ ,  $O(N)$ , ...
- Compute Complexity
  - $O(1)$ ,  $O(\log N)$ ,  $O(N)$ , ...
- Compute Consistency
  - Uniform vs. non-uniform
- Total / Partial
  - Complete vs. sparse
- One-to-one / Many-to-one
  - Uniform vs. adaptive



# Data Structure Examples

- Brook streams (Buck et al. 2004)



1D Virtual

2D Physical



Aaron Lefohn  
University of California, Davis

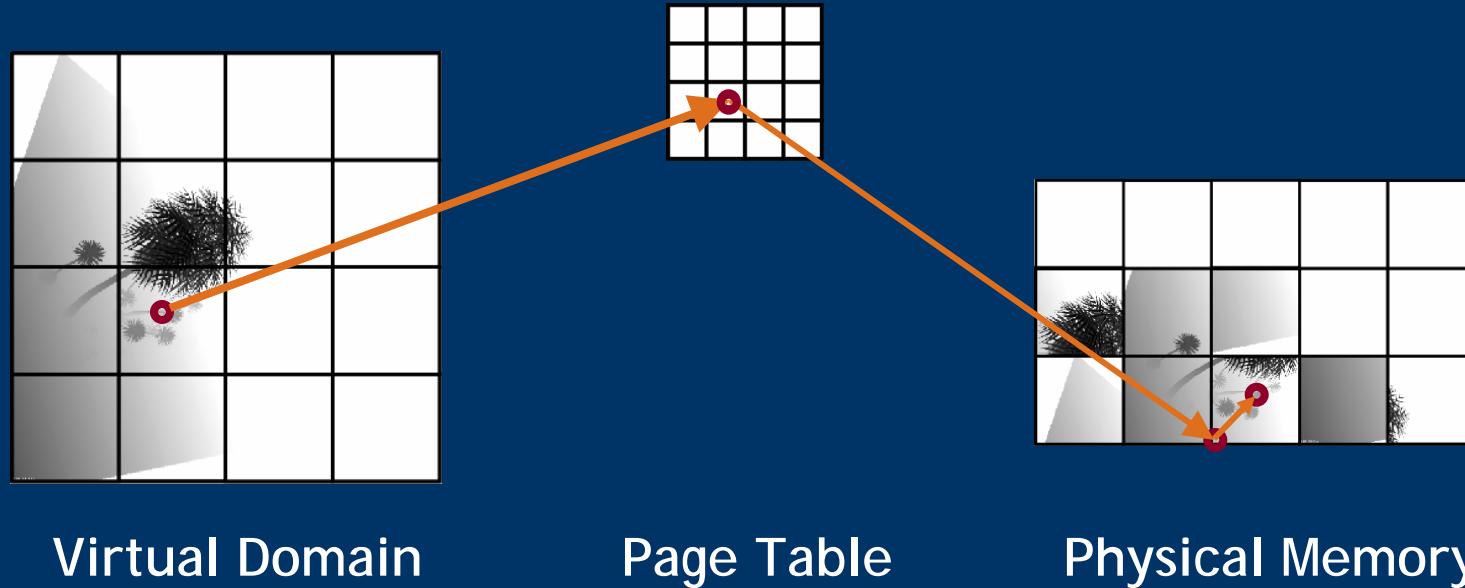
# Data Structure Examples

- Brook streams (Buck et al. 2004)
  - Physical address 2D
  - Virtual address N-D
  - Address translator
    - Analytic
    - $O(1)$  memory
    - $O(1)$  compute
    - Uniform consistency
    - Total, uniform mapping



# Data Structure Examples

- Dynamic sparse 3D grid (Lefohn et al. 2003)



Virtual Domain

Page Table

Physical Memory



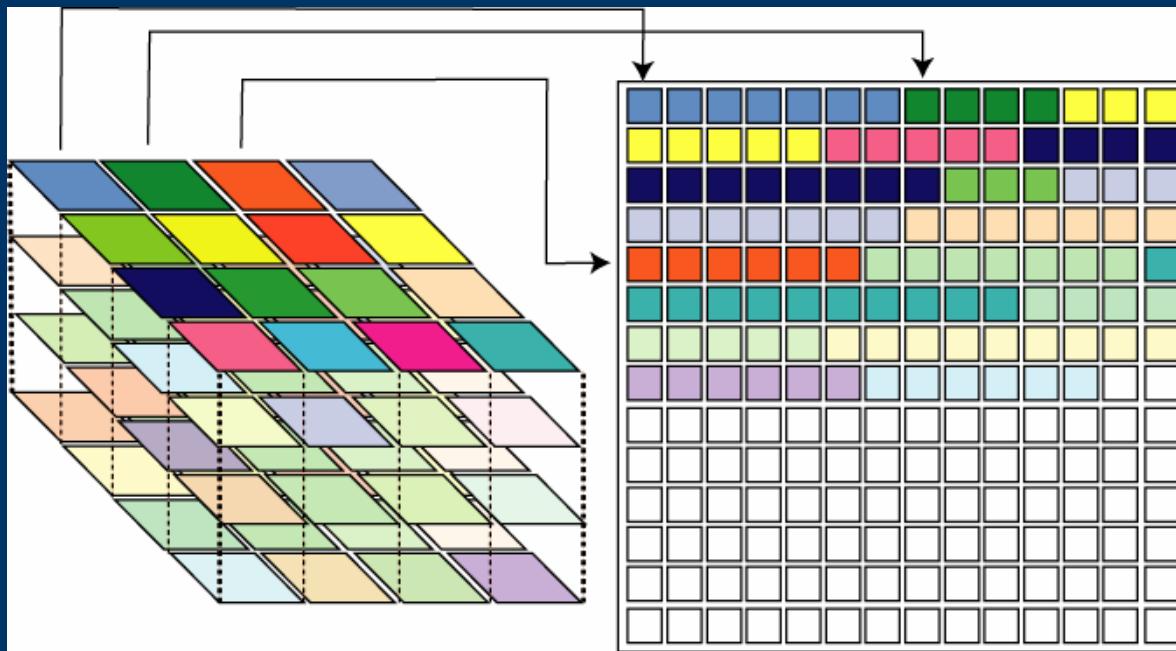
# Data Structure Examples

- Dynamic sparse 3D grid (Lefohn et al. 2003)
  - Physical address 2D
  - Virtual address 3D
  - Address translator
    - Discrete
    - $O(N)$  memory
    - $O(1)$  compute
    - Uniform consistency
    - Partial, uniform mapping



# Data Structure Examples

- Photon Map (kNN-grid) (Purcell et al. 2003)



*Image from "Implementing Efficient Parallel Data Structures on GPUs,"  
Lefohn et al., GPU Gems II, ch. 33, 2005*



# Data Structure Examples

- Photon Map (kNN-grid) (Purcell et al. 2003)
  - Physical address 2D
  - Virtual address 3D
  - Address translator 3D page table
    - Variable sized phys pages
    - "Grid of lists"
    - Discrete
    - $O(N)$  memory
    - $O(L)$  compute
    - Non-uniform consistency
    - Partial, adaptive mapping



# Glift Iterators

---

- We've so far only discussed data *access*
- What about data structure *traversal*?



# Iterators

---

- Separate algorithms and data structures
  - Minimal interface between data and algorithm
  - Required for GPGPU use of data structure
  - Encapsulate GPGPU optimizations



# Iterators

- Abstract data access and traversal

```
DataStructureType::iterator it;  
for (it = data.begin(); it != data.end(); ++it)  
{  
    *it = -(*it);  
}
```



# Glift Iterators

---

- **Address iterators**
  - Iterator value is N-D address
  - GPU interpolants
- **Element iterators**
  - Iterator value is data structure element
  - C/C++ pointer, STL iterator, streams



# Element Iterator Concepts

---

- **Permission**
  - Read-only, write-only, read-write
- **Access region**
  - Single, neighborhood, random
- **Traversal**
  - Forward, backward, parallel range



# Which Element Iterators?

---

- Read-only, single access, range iterator
  - GPU stream input
- Read-only, random-access, range iterator
  - GPU texture input
- Write-only, single access, range iterator
  - GPU render target



# Example 1 : “Before” and “After” Glift

---

- Transform GPU code with Glift



Aaron Lefohn  
University of California, Davis

# Simple Example

- 3D Array with 2D physical memory

CPU (C++)

```
float srcData[10][10][10];
float dstData[10][10][10];

... initialize data ...

for (size_t z = 1; z < 10; ++z) {
    for (size_t y = 1; z < 10; ++y) {
        for (size_t x = 1; z < 10; ++x) {
            dstData[z][y][x] = srcData[z-1][y-1][x-1];
        }
    }
}
```



# Example 1: Shader w/out Glift

```
float3 physToVirt( float2 pa, float2 physSize, float3 virtSizes ) {
    float3 va;
    float addr1D = pa.y * physSize.x + pa.x;
    va.z = floor( addr1D / virtSizes.z );
    addr1D -= va.z * sizeConst3D.z;
    va.y = floor( addr1D / virtSizes.y );
    va.x = addr1D - va.y * virtSizes.y;
    return va;
}
```

Physical-to-Virtual  
Address Translation

```
float2 virtToPhys( float3 va, float2 physSize, float3 virtSizes ) {
    float addr1D = dot( va, virtSizes );
    float normAddr1D = addr1D / physSize.x;
    float2 pa = float2(frac(normAddr1D) * physSize.x, normAddr1D);
}
```

Virtual-to-Physical  
Address Translation

```
float3 main( uniform samplerRECT physMem,
              uniform float2          physSize,
              uniform float3          virtSizes,
              uniform float2          pa : WPOS ) : COLOR
{
    float3 va = physToVirt( floor(pa), physSize, virtSizes );
    float3 neighborAddr = va - float3(1, 1, 1);
    return texRECT(data, virtToPhys(neighborAddr3D, physSize, virtSizes));
}
```

Physical Memory Read



# Example 1: Glift Components

```
float3 physToVirt( float2 pa, float2 physSize, float3 virtSizes ) {
    float3 va;
    float addr1D = pa.y * physSize.x + pa.x;
    va.z = floor( addr1D / virtSizes.z );
    addr1D -= va.z * sizeConst3D.z;
    va.y = floor( addr1D / virtSizes.y );
    va.x = addr1D - va.y * virtSizes.y;
    return va;
}
```

Address Iterator

```
float2 virtToPhys( float3 va, float2 physSize, float3 virtSizes ) {
    float addr1D = dot( va, virtSizes );
    float normAddr1D = addr1D / physSize.x;
    float2 pa = float2(frac(normAddr1D) * physSize.x, normAddr1D);
}
```

VirtMem

```
float3 main( uniform samplerRECT physMem,
              uniform float2 physSize,
              uniform float3 virtSizes,
              float2 pa : WPOS ) : COLOR
{
    float3 va = physToVirt( floor(pa), physSize, virtSizes );
    float3 neighborAddr = va - float3(1, 1, 1);
    return texRECT(data, virtToPhys(neighborAddr, physSize, virtSizes) );
}
```

VirtMem



# Example 1: GPU Shader with Glift

## Cg Usage

```
float3 main( uniform VMem3D srcData,
              AddrIter3D      iter ) : COLOR
{
    float3 va = iter.value();
    return srcData.vTex3D( va - float3(1,1,1) );
}
```



# Example 1: Glift Data Structures

## C++ Usage

```
vec3i origin(0,0,0);
vec3i size(10,10,10);

typedef ArrayGpu<vec3i,vec1f> ArrayType;
ArrayType srcData( size );
ArrayType dstData( size );

... initialize dataPtr ...
srcData.write( origin, size, dataPtr );

typedef ArrayType::addr_trans AddrTransType;
AddrTransType::gpu_range it =
    dstData.addr_trans().gpu_range(origin, size);

it.bind_for_read( iterCgParam );
srcData.bind_for_read( srcCgParam );
dstData.bind_for_write( COLOR0, myFrameBufferObject );

exec_gpu_iterators( it );
```



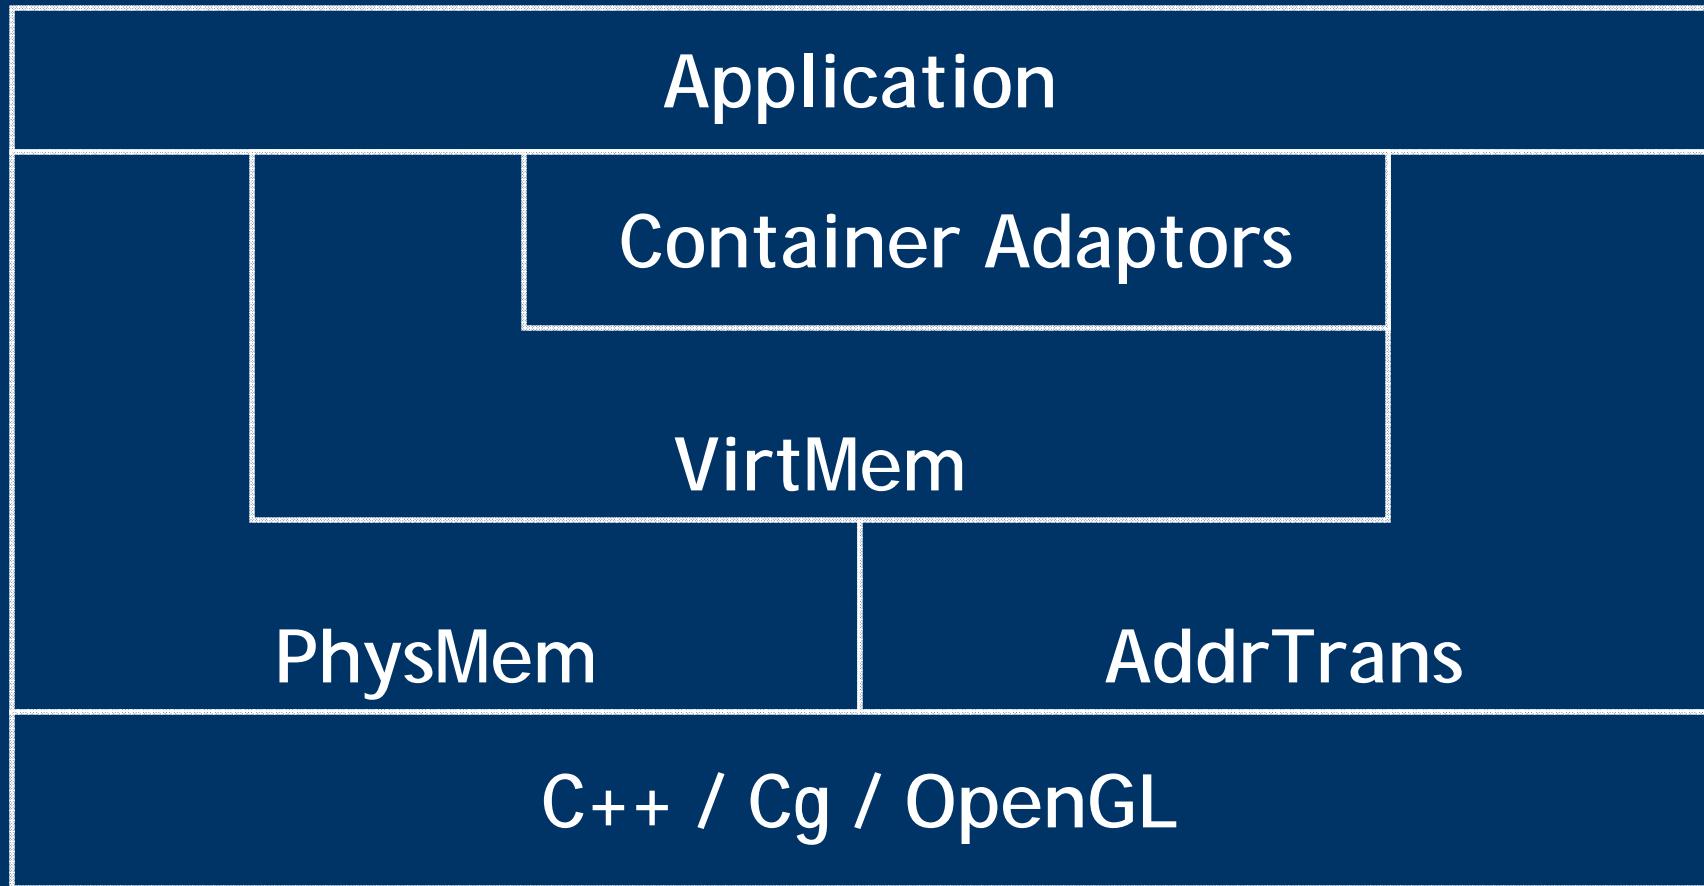
# Overview

---

- Motivation
- Abstraction
- Implementation
- Examples
- Conclusions



# Glift Components



# Glift Design Goals

---

- Efficiency
- Easy, incremental adoption
- Easily extensible
- CPU/GPU interoperability



# Glift Design Goals

---

- **Efficiency**
  - Static polymorphism (C++ and Cg)
  - Cg program specialization
  - Cg compiler optimizations
- **Easy, incremental adoption**
- **Easily extensible**
- **CPU/GPU interoperability**



# Glift Design Goals

---

- Efficiency
- Easy, incremental adoption
  - Integrate with Cg/OpenGL/C++
  - STL-like and texture-like interfaces
  - Use components alone or composited
- Easily extensible
- CPU/GPU interoperability



# Glift Design Goals

---

- Efficiency
- Easy, incremental adoption
- Easily extensible
  - Create new structure by:
    - Change behavior of existing address translator
    - New address translator
    - New container adaptor
- CPU/GPU interoperability



# Glift Design Goals

---

- Efficiency
- Easy, incremental adoption
- Easily extensible
- CPU/GPU interoperability
  - Unified C++/Cg code base
  - Map memory to CPU or GPU
  - CPU and GPU iterators



# C++/Cg Integration

- Each component defines C++ and Cg code
  - C++ objects have Cg struct representation
  - Stringified Cg parameterized by C++ templates
- Cg “template” instantiation
  - Insert generated Glift source code into shader

```
glift::cgGetTemplateType<MyDataStructType>();
glift::cgInstantiateParameter(...);
```

- All other compilation/loading/binding identical to standard shader



# Cg Compilation Example

- Cg code

```
float4 main( uniform VMem3D octree,
              float3 coord ) : COLOR
{
    return octree.vMem3D(coord);
}
```

- C++ code

```
typedef OctreeGPU<vec4ub> octree_type;
GliftType type = cgGetTemplateType<octree_type>();
CGprogram prog = cgCreateProgram(...);
prog = cgInstantiateParameter(prog, "octree", type);
cgCompileProgram(prog);
```



# Overview

---

- Motivation and previous work
- Abstraction
- Case Study
  - Adaptive shadow maps and octree 3D paint
- Conclusions



# Example 2: Adaptive Shadow Maps

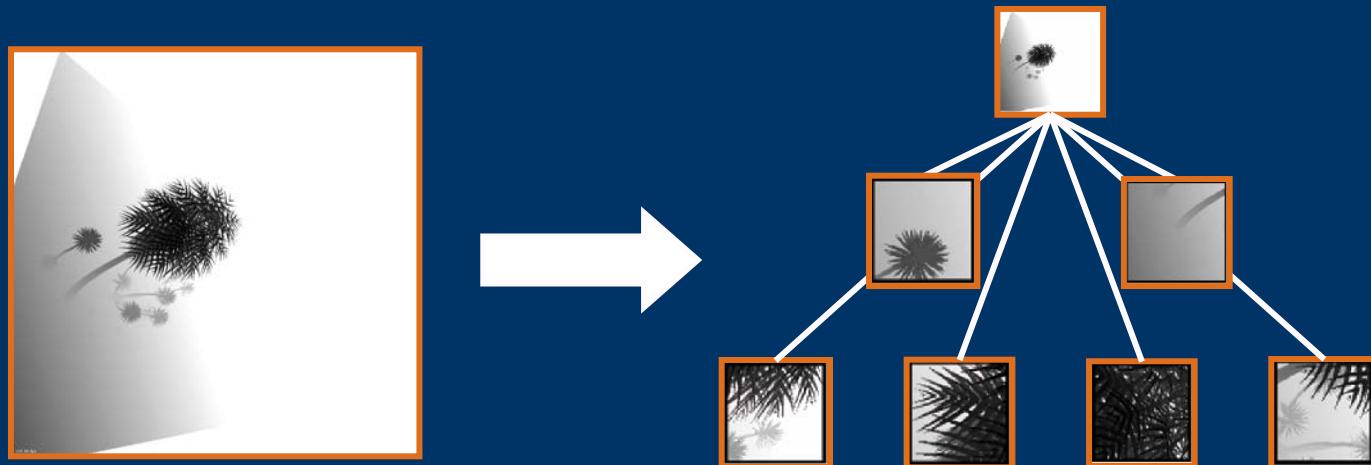
---

- Show Glift usage with
  - Complex application
  - Complex data structure



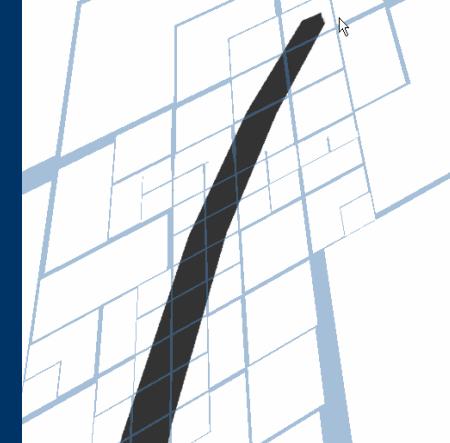
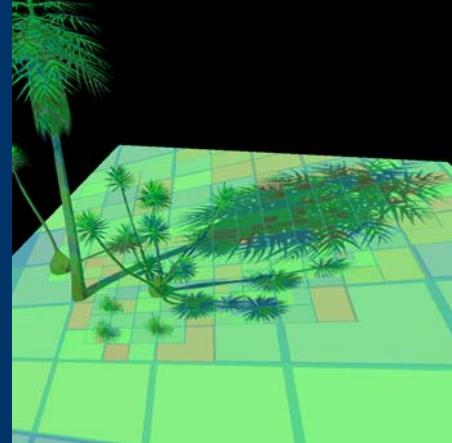
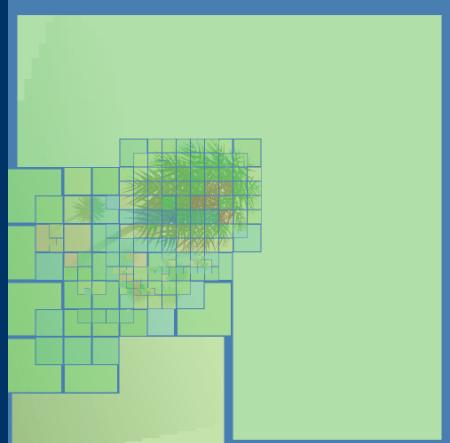
# Example 2: Adaptive Shadow Maps

- Fernando et al., ACM SIGGRAPH 2001
- Elegant solution to shadow map aliasing
  - Quadtree of small shadow maps
  - Shadow maps need resolution only on shadow boundary
  - Required resolution determined by projected area of screen space pixel into light space



# Adaptive Shadow Maps

- Why Adaptive Shadow Maps with Glift?
  - Many recent (2004) shadow papers cite ASMs as high quality solution but not possible on graphics hardware
  - Algorithm is simple. Data structure is hard.



# Adaptive Shadow Map Algorithm

---

- **Iterative refinement algorithm**
  - Identify shadow pixels w/ resolution mismatch
  - Create small shadow map “pages” at requested resolution
- **Shadow lookup**
  - Compute shadow map coordinate and resolution
  - Lookup in ASM (tree of small shadow map pages)
- **ASM depends on both camera and light position!**



# ASM Data Structure Requirements

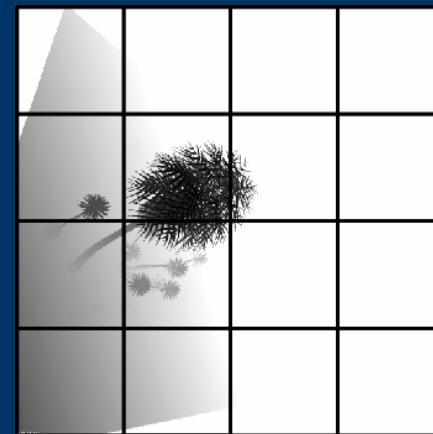
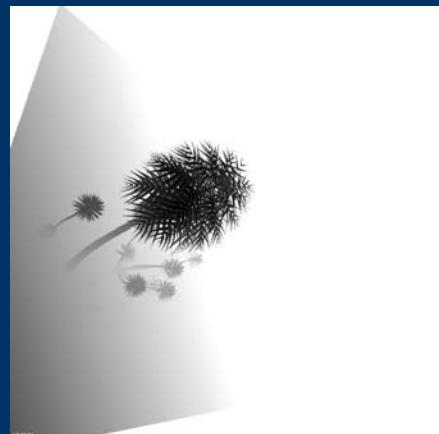
---

- Adaptive
- Multiresolution
- Fast, parallel random-access read
  - 2x2 native Percentage Closer Filtering (PCF)
  - Trilinear interpolated mipmapped PCF
- Fast, parallel write
- Fast, parallel insert and erase



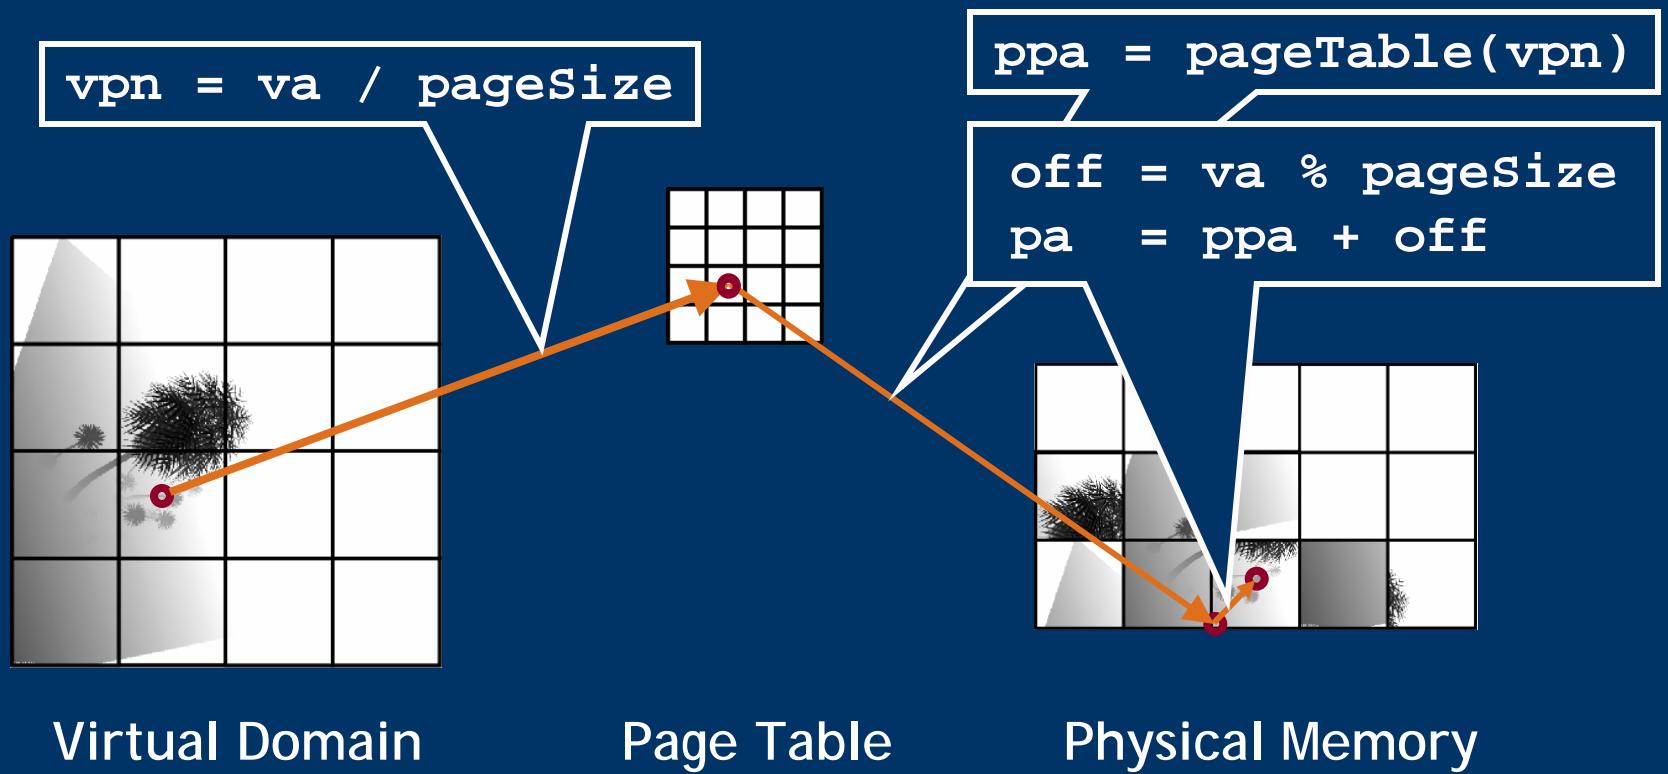
# ASM Data Structure

- Start with page table address translator
  - Coarse, uniform discretization of virtual domain
    - $O(N)$  memory                       $O(1)$  insert
    - $O(1)$  computation                   $O(1)$  erase
    - Uniform consistency
    - Partial mapping (sparse)



# ASM Data Structure

- Page table example



# ASM Data Structure Requirements

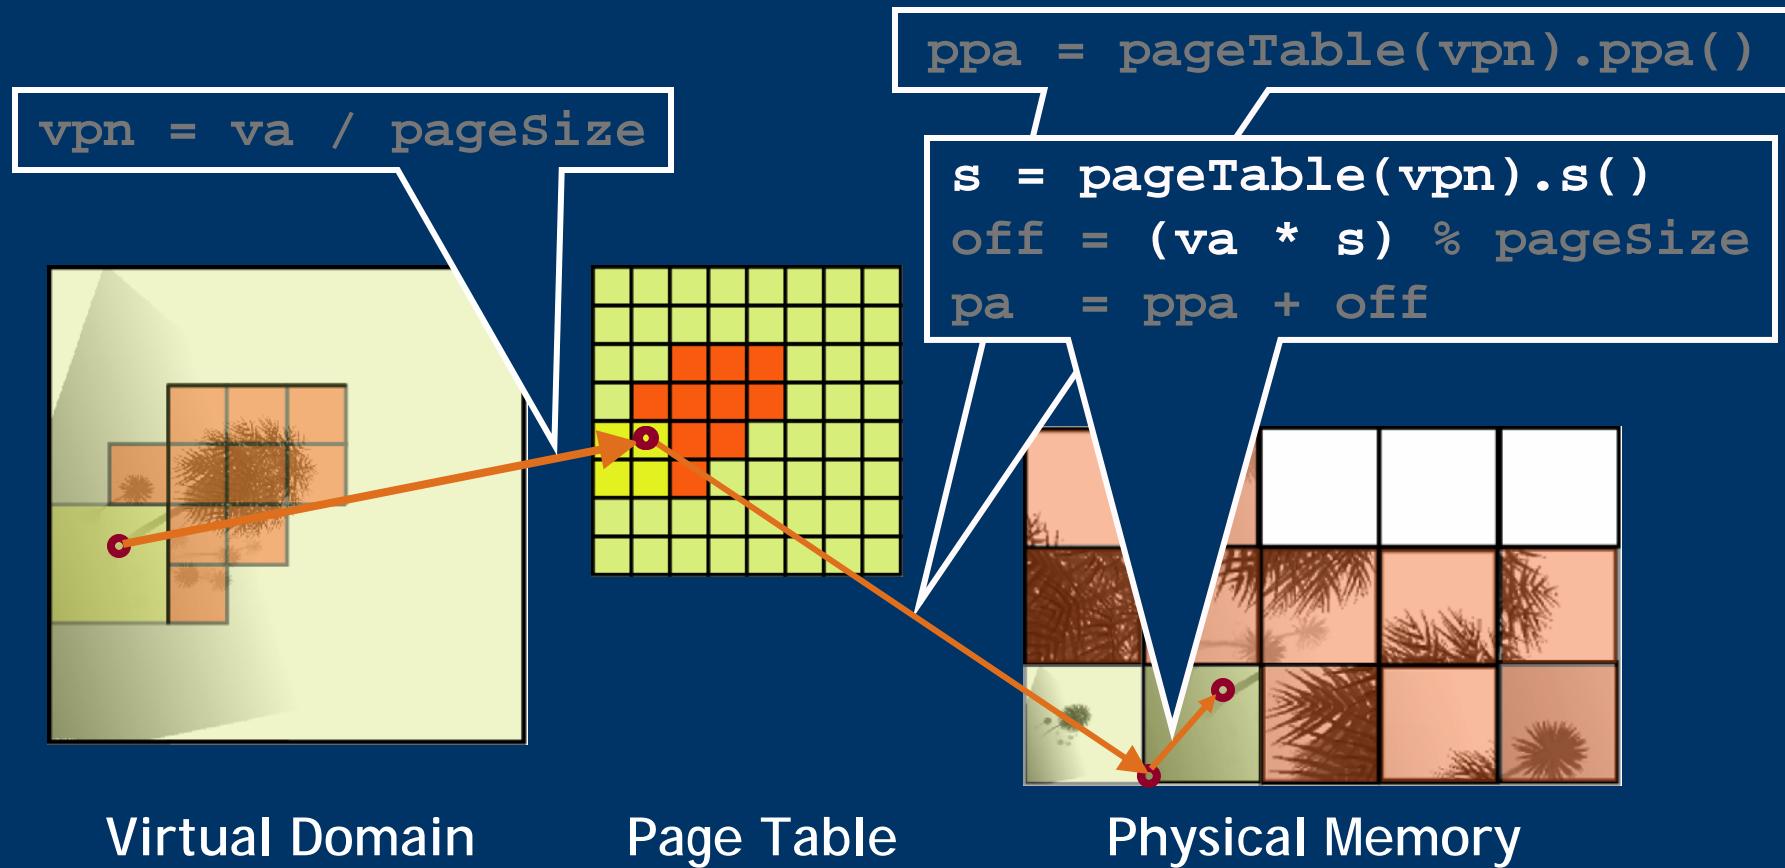
---

- Adaptive
- Multiresolution
- Fast, parallel random-access read
  - 2x2 native Percentage Closer Filtering (PCF)
  - Trilinear interpolated mipmapped PCF
- Fast, parallel write
- Fast, parallel insert and erase



# ASM Data Structure

- Adaptive Page Table
  - Map multiple virtual pages to single physical page



# ASM Data Structure Requirements

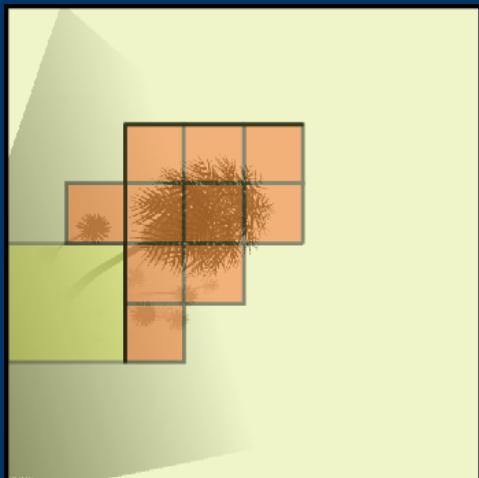
---

- Adaptive
- Multiresolution
- Fast, parallel random-access read
  - 2x2 native Percentage Closer Filtering (PCF)
  - Trilinear interpolated mipmapped PCF
- Fast, parallel write
- Fast, parallel insert and erase

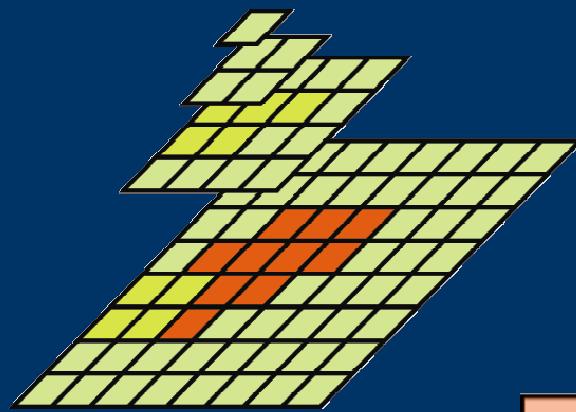


# ASM Data Structure

- Multiresolution Page Table



Virtual Domain



Mipmap  
Page Table



Physical Memory



# ASM Data Structure Requirements

---

- Adaptive
- Multiresolution
- Fast, parallel random-access read
  - 2x2 native Percentage Closer Filtering (PCF)
  - Trilinear interpolated mipmapped PCF
- Fast, parallel write
- Fast, parallel insert and erase



# ASM Data Structure Requirements

---

- How support bilinear filtering?
  - Duplicate 1 column and 1 row of texels in each page
- Mipmapped trilinear?
  - “By-hand” interpolation between mipmap levels



# ASM Data Structure Requirements

---

- Adaptive
- Multiresolution
- Fast, parallel random-access read
  - 2x2 native Percentage Closer Filtering (PCF)
  - Trilinear interpolated mipmapped PCF
- Fast, parallel write
- Fast, parallel insert and erase



# How Define ASM Structure in Glift?

- Start with generic page table AddrTrans
  - Use mipmapped PhysMem for page table
  - Change template parameter to add adaptivity
- Write page allocator
  - `alloc_pages, free_pages`
- Finally...

```
typedef PageTableAddrTrans<...>           PageTable;
typedef PhysMemGPU<vec2f, vec1s>             PMem2D;
typedef VirtMemGPU<PageTable, PMem2D>         VPageTable;
typedef AdaptiveMem<VPageTable, PageAllocator> ASM;
```



# ASM Data Structure Usage

```
float4 main( uniform VMem2D asm,  
             float3 shadowCoord,  
             float4 litColor      ) : COLOR  
{  
    float isInLight = asm.vTex2Ds( shadowCoord );  
    return lerp( black, litColor, isInLight );  
}  
  
asm.bind_for_read( ... );  
asm.bind_for_write( ... );  
  
asm.alloc_pages( ... );  
asm.free_page( ... );  
  
...
```

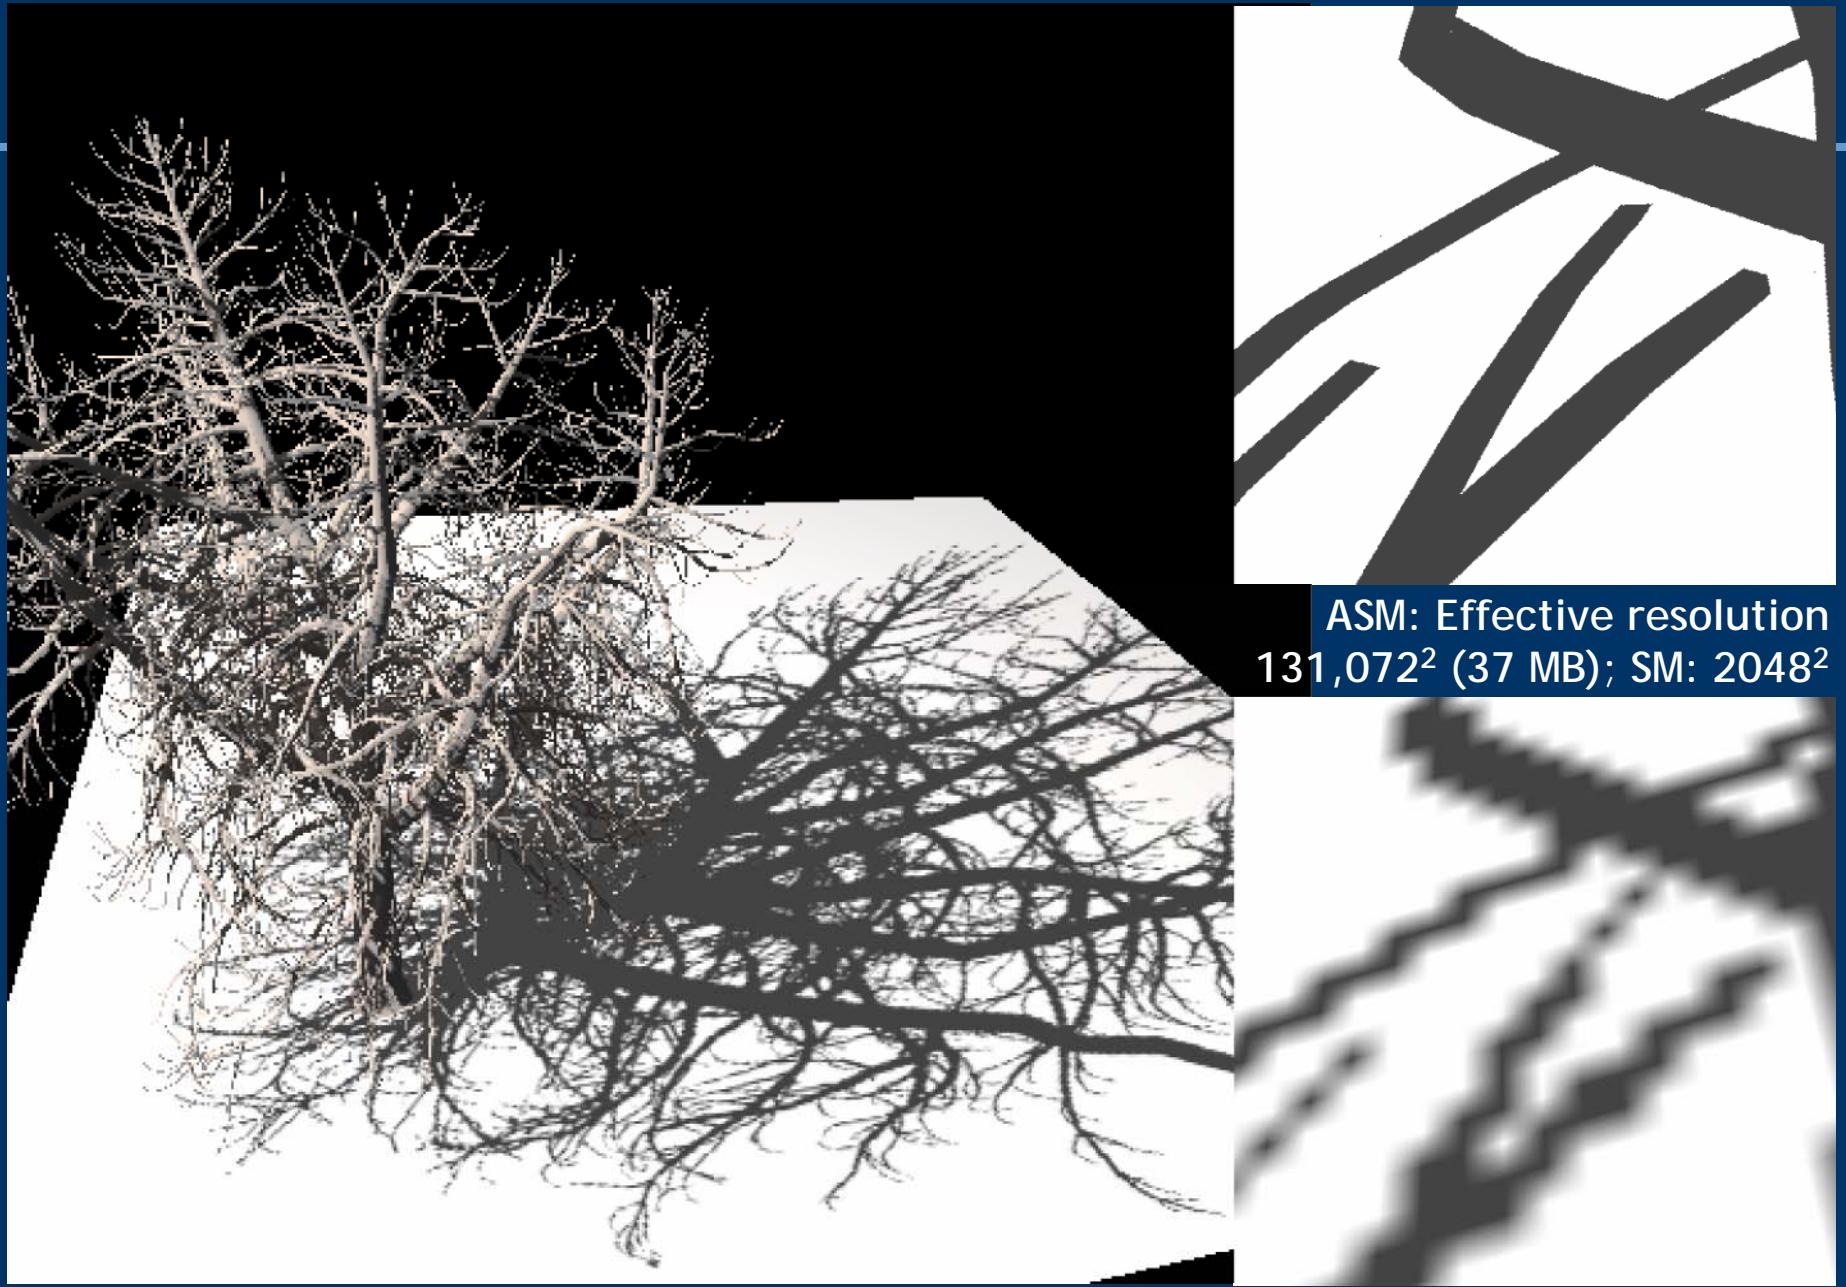


# Adaptive Shadow Map Algorithm

---

- Faithful to Fernando et al. 2001
- Refinement algorithm
  - Identify shadow pixels w/ resolution mismatch (GPU)
  - Compact pixels into small stream (GPU)
  - CPU reads back compacted stream (GPU→CPU)
  - Allocate pages
    - Draw new PTEs into mipmap page tables (CPU→GPU)
    - Draw depth into ASM for each new page (GPU)





Aaron Lefohn  
University of California, Davis

[Thanks to Yong Kil for the tree model]

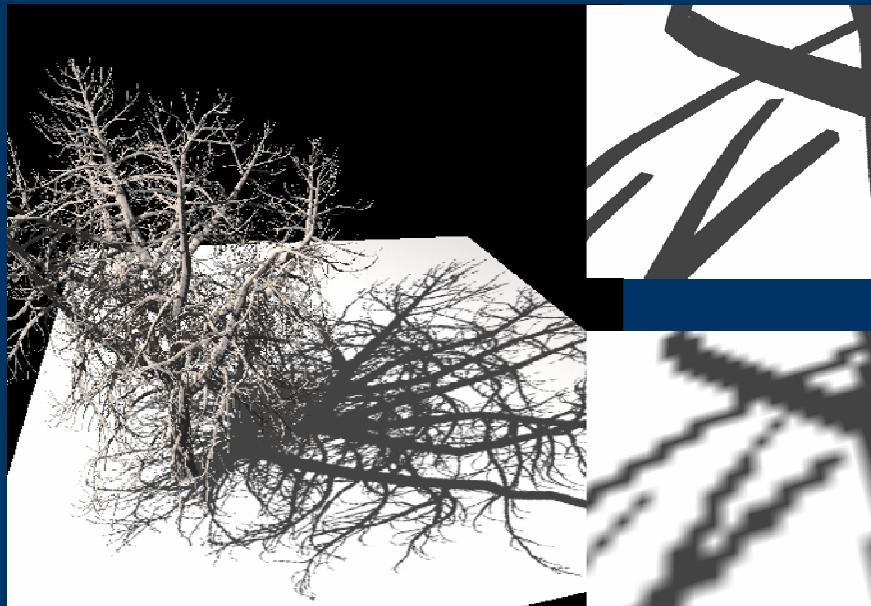
# “Octree” 3D Paint

- Interactive painting on unparameterized 3D surfaces
- 3D version of ASM data structure
- Differs from previous work:
  - Quadrilinear filtering
  - $O(1)$ , uniform access
- Interactive with effective resolutions between  $64^3$  and  $2048^3$



# Demo

---



Aaron Lefohn  
University of California, Davis

# ASM Results

---

- Effective shadow map resolution up to  $131,072^2$ 
  - 16<sup>2</sup> - 64<sup>2</sup> page size
  - 512<sup>2</sup> - 2048<sup>2</sup> page table
  - 2048<sup>2</sup> - 4096<sup>2</sup> physical memory
  - 20 - 80 MB
- Performance (45k polygon model)
  - 15 fps while moving camera (including refinement)
  - 5-10 fps while moving light
- Lookup time compared to 2048<sup>2</sup> shadow map:
  - Bilinear filtered: 90% performance of traditional
  - Trilinear filtered mipmapped: 73%



# Glift Results

- Static instruction results
  - With Cg program specialization

|                 | Glift | By-Hand | Brook |
|-----------------|-------|---------|-------|
| • 1D → 2D       | 4     | 3       | 4     |
| • 3D page table | 5     | 5       |       |
| • ASM           | 9     | 9       |       |
| • Octree        | 10    | 9       |       |
| • ASM + offset  | 10    | 9       |       |

- Conclusion : Glift structures within 1 instr of hand-coded Cg

Measured with NVShaderPerf, NVIDIA driver 75.22, Cg 1.4a



# Overview

---

- Motivation and previous work
- Abstraction
- Implementation
- Examples
- Conclusions



# Summary

---

- GPU programming needs data structure abstraction
  - Separate data structures and algorithms
  - More complex data structures and algorithms
- Why programmable address translation?
  - Common pattern in GPU data structures
  - Small amount of code virtualizes GPU memory model



# Summary

---

- **Glift template library**
  - Generic C++/Cg implementation of abstraction
  - Nearly as efficient as hand coding
  - Integrates with OpenGL/Cg
- **Iterator computation model**
  - Generalize GPU computation model
  - Can future rasterizer increment iterators?



# Acknowledgements

- Craig Kolb, Nick Triantos, Cass Everitt NVIDIA
  - Fabio Pellacini Dartmouth
  - Adam Moerschell, Yong Kil UCDavis  
Serban Porumbescu, Chris Co, ....
  - Ross Whitaker, Chuck Hansen, Milan Ikits U. of Utah
  - National Science Foundation Graduate Fellowship
  - Department of Energy



# More Information

---

- Paper in ACM Transactions on Graphics (Jan. 2006)
  - “Glift : Generic, Efficient, Random-Access GPU Data Structures”
- ACM SIGGRAPH 2005 Sketches
  - “Dynamic Adaptive Shadow Maps on Graphics Hardware”
  - “Octree Texture on Graphics Hardware”
- Google “Glift”
  - <http://graphics.cs.ucdavis.edu/~lefohn/>

